

Book Reader Testing PDF

Purpose:

This PDF is generated specifically for testing TOC extraction and chapter processing.

Contents:

- a visual table of contents page
- bookmark-based PDF TOC entries
- five top-level chapters so the current chapter segmentation code yields four usable chapters
- repeated, predictable text for stable debugging

Table of Contents

Chapter 1: Foundations of Reading Systems	3
1.1 What a Book Pipeline Does	3
1.2 Pages, Text, and Structure	3
Chapter 2: Extracting a Table of Contents	4
2.1 Identifying Heading Levels	4
2.2 Mapping Headings to Pages	4
Chapter 3: Segmenting Chapters	5
3.1 Splitting Pages by Headings	5
3.2 Carrying Context Across Pages	5
Chapter 4: Generating Study Slides	6
4.1 Summaries for Each Topic	6
4.2 Turning Notes into Slides	6
Chapter 5: Asking Questions About Content	7
5.1 Retrieving Relevant Notes	7
5.2 Producing Helpful Answers	7

Chapter 1: Foundations of Reading Systems

Page 3

1.1 What a Book Pipeline Does

This section belongs to chapter 1 of the generated testing book.

It exists to exercise TOC extraction, heading matching, and chapter segmentation.

The text is intentionally plain so PDF text extraction stays stable across runs.

Each heading appears directly in the page body to help similarity-based matching.

Use this file when debugging page slicing, slide generation, and question answering.

Key points:

- headings should match the bookmark titles
- chapter pages should be easy to inspect manually
- extracted text should remain simple and predictable

Example note:

A reliable testing PDF is more useful than a realistic-looking but inconsistent one.

1.2 Pages, Text, and Structure

This section belongs to chapter 1 of the generated testing book.

It exists to exercise TOC extraction, heading matching, and chapter segmentation.

The text is intentionally plain so PDF text extraction stays stable across runs.

Each heading appears directly in the page body to help similarity-based matching.

Use this file when debugging page slicing, slide generation, and question answering.

Key points:

- headings should match the bookmark titles
- chapter pages should be easy to inspect manually
- extracted text should remain simple and predictable

Example note:

A reliable testing PDF is more useful than a realistic-looking but inconsistent one.

Chapter 2: Extracting a Table of Contents

Page 4

2.1 Identifying Heading Levels

This section belongs to chapter 2 of the generated testing book.

It exists to exercise TOC extraction, heading matching, and chapter segmentation.

The text is intentionally plain so PDF text extraction stays stable across runs.

Each heading appears directly in the page body to help similarity-based matching.

Use this file when debugging page slicing, slide generation, and question answering.

Key points:

- headings should match the bookmark titles
- chapter pages should be easy to inspect manually
- extracted text should remain simple and predictable

Example note:

A reliable testing PDF is more useful than a realistic-looking but inconsistent one.

2.2 Mapping Headings to Pages

This section belongs to chapter 2 of the generated testing book.

It exists to exercise TOC extraction, heading matching, and chapter segmentation.

The text is intentionally plain so PDF text extraction stays stable across runs.

Each heading appears directly in the page body to help similarity-based matching.

Use this file when debugging page slicing, slide generation, and question answering.

Key points:

- headings should match the bookmark titles
- chapter pages should be easy to inspect manually
- extracted text should remain simple and predictable

Example note:

A reliable testing PDF is more useful than a realistic-looking but inconsistent one.

Chapter 3: Segmenting Chapters

Page 5

3.1 Splitting Pages by Headings

This section belongs to chapter 3 of the generated testing book.
It exists to exercise TOC extraction, heading matching, and chapter segmentation.
The text is intentionally plain so PDF text extraction stays stable across runs.
Each heading appears directly in the page body to help similarity-based matching.
Use this file when debugging page slicing, slide generation, and question answering.

Key points:

- headings should match the bookmark titles
- chapter pages should be easy to inspect manually
- extracted text should remain simple and predictable

Example note:

A reliable testing PDF is more useful than a realistic-looking but inconsistent one.

3.2 Carrying Context Across Pages

This section belongs to chapter 3 of the generated testing book.
It exists to exercise TOC extraction, heading matching, and chapter segmentation.
The text is intentionally plain so PDF text extraction stays stable across runs.
Each heading appears directly in the page body to help similarity-based matching.
Use this file when debugging page slicing, slide generation, and question answering.

Key points:

- headings should match the bookmark titles
- chapter pages should be easy to inspect manually
- extracted text should remain simple and predictable

Example note:

A reliable testing PDF is more useful than a realistic-looking but inconsistent one.

Chapter 4: Generating Study Slides

Page 6

4.1 Summaries for Each Topic

This section belongs to chapter 4 of the generated testing book.

It exists to exercise TOC extraction, heading matching, and chapter segmentation.

The text is intentionally plain so PDF text extraction stays stable across runs.

Each heading appears directly in the page body to help similarity-based matching.

Use this file when debugging page slicing, slide generation, and question answering.

Key points:

- headings should match the bookmark titles
- chapter pages should be easy to inspect manually
- extracted text should remain simple and predictable

Example note:

A reliable testing PDF is more useful than a realistic-looking but inconsistent one.

4.2 Turning Notes into Slides

This section belongs to chapter 4 of the generated testing book.

It exists to exercise TOC extraction, heading matching, and chapter segmentation.

The text is intentionally plain so PDF text extraction stays stable across runs.

Each heading appears directly in the page body to help similarity-based matching.

Use this file when debugging page slicing, slide generation, and question answering.

Key points:

- headings should match the bookmark titles
- chapter pages should be easy to inspect manually
- extracted text should remain simple and predictable

Example note:

A reliable testing PDF is more useful than a realistic-looking but inconsistent one.

Chapter 5: Asking Questions About Content

Page 7

5.1 Retrieving Relevant Notes

This section belongs to chapter 5 of the generated testing book.

It exists to exercise TOC extraction, heading matching, and chapter segmentation.

The text is intentionally plain so PDF text extraction stays stable across runs.

Each heading appears directly in the page body to help similarity-based matching.

Use this file when debugging page slicing, slide generation, and question answering.

Key points:

- headings should match the bookmark titles
- chapter pages should be easy to inspect manually
- extracted text should remain simple and predictable

Example note:

A reliable testing PDF is more useful than a realistic-looking but inconsistent one.

5.2 Producing Helpful Answers

This section belongs to chapter 5 of the generated testing book.

It exists to exercise TOC extraction, heading matching, and chapter segmentation.

The text is intentionally plain so PDF text extraction stays stable across runs.

Each heading appears directly in the page body to help similarity-based matching.

Use this file when debugging page slicing, slide generation, and question answering.

Key points:

- headings should match the bookmark titles
- chapter pages should be easy to inspect manually
- extracted text should remain simple and predictable

Example note:

A reliable testing PDF is more useful than a realistic-looking but inconsistent one.